

## STAT 656 Homework 3

E. Lee Rainwater – lee.rainwater@tamu.edu

### 1 Executive Summary

The supplied data, from *credithistory.xlsx*, contained 1000 observations, the details of which are listed in Appendix 1. The response variable was taken to be the feature, *good\_bad*, which is a binary classification of whether the applicant is expected to have good or bad credit.

A total of ten logistic regression models were run:

- Logistic with all features with model validation
- Logistic with stepwise feature selection
- Logistic with regularization
  - Regularization coefficients ranging from  $1.0 \times 10^{-4}$  to  $+\infty$

For each model, the dataset was randomly split 70%/30% into test/validation datasets. Metrics for each of the validation cases are shown in Table 1:

Model	Precision	Sensitivity	Validation		
			Specificity	F1	MISC
All Features with Validation	0.8186	0.8739	0.4487	0.8453	23.7%
Logistic with Stepwise	0.8238	0.9054	0.4487	0.8627	21.3%
Regularization Logistic, C = 0.0001	0.7492	0.9955		0.8549	25.0%
Regularization Logistic, C = 0.01	0.7762	0.9685		0.8617	23.0%
Regularization Logistic, C = 0.1	0.8118	0.9324		0.8679	21.0%
Regularization Logistic, C = 1	0.8139	0.8468		0.8300	25.7%
Regularization Logistic, C = 5	0.8162	0.8604		0.8377	24.7%
Regularization Logistic, C = 10	0.8194	0.8378		0.8285	25.7%
Regularization Logistic, C = 50	0.8178	0.8694		0.8428	24.0%
Regularization Logistic, C = $\infty$	0.8248	0.8694		0.8465	23.3%

Table 1 - Model Validation Metrics

On cursory investigation, the models appear to be very close in quality. Given that the target variable is the binary determination of the customer's credit (good/bad), it may be surmised that the effects are of false negatives vs. false positives. A false negative will result in a loan being not granted to a qualified customer, resulting in the loss of associated profit. A false positive, on the other hand, *may* result in the loss of the entire value of the loan, plus expenses. Thus, it is surmised that the negative consequences of a false positive are greater than those of a false negative.

For this reason, *sensitivity* is chosen as the primary metric for evaluation, as it more strongly considers the effect of false positives. As can be seen in Table 1, *Regularization Logistic, C = 0.0001* provides the highest degree of sensitivity.

A more in-depth analysis would include consideration of the actual amount of profit loss due to false negatives, as the lower *precision* value for this same model indicates that more applicants may be falsely rejected.

## 2 Appendix 1 – Python Output

### 2.1 Imputation & Outliers

```
..... Missing Outliers
age..... 35 6
amount.... 12 9
checking.. 0 0
coapp..... 12 0
depends... 0 0
duration.. 42 0
employed.. 0 6
existcr... 0 0
foreign... 0 0
history... 0 0
housing... 0 0
installp.. 0 0
job..... 0 0
marital... 9 5
other..... 0 0
property.. 0 0
purpose... 564 0
resident.. 11 0
savings... 4 2
telephon.. 19 0
good_bad.. 0 0
```

### 2.2 Logistic Regression – All Features with Validation

Python 3.7.6 (default, Jan 8 2020, 13:42:34)  
Type "copyright", "credits" or "license" for more information.

IPython 7.13.0 -- An enhanced Interactive Python.

```
runfile('/Users/edwardrainwater/OneDrive - Texas A&M University/Summer-2020/STAT 656
Applied Analytics/hw-03/rainwater-stat656-hw03.py', wdir='/Users/edwardrainwater/OneDrive
- Texas A&M University/Summer-2020/STAT 656 Applied Analytics/hw-03')
```

	checking	duration	history	purpose	...	depends	telephon	foreign	good_bad
0	1	NaN	4	3	...	1	NaN	1	good
1	2	48.0	2	NaN	...	1	NaN	1	bad
2	4	12.0	4	NaN	...	2	NaN	1	good
3	1	42.0	2	NaN	...	2	NaN	1	good
4	1	24.0	3	NaN	...	2	NaN	1	bad

[5 rows x 21 columns] (1000, 21)

```
checking      int64
duration      float64
history       int64
purpose       object
amount        float64
savings       float64
employed      int64
installp      int64
marital       float64
```

```

coapp      float64
resident   float64
property    int64
age         float64
other       int64
housing     int64
existcr     int64
job         int64
depends      int64
telephon    float64
foreign     int64
good_bad    object
dtype: object

```

\*\*\*\*\* Data Preprocessing \*\*\*\*\*

Features Dictionary Contains:

3 Interval,  
4 Binary,  
13 Nominal, and  
1 Excluded Attribute(s).

Data contains 1000 observations & 21 columns.

/Users/edwardrainwater/opt/anaconda3/envs/Stat656MacOSX/lib/python3.7/site-packages/AdvancedAnalytics/ReplaceImputeEncode.py:338: UserWarning: purpose:has more than 50% missing.Recommend setting Data Type set to DT.Ignore.

"Recommend setting Data Type set to DT.Ignore.")

#### Attribute Counts

.....	Missing	Outliers
age.....	35	6
amount....	12	9
checking..	0	0
coapp.....	12	0
depends...	0	0
duration..	42	0
employed..	0	6
existcr...	0	0
foreign...	0	0
history...	0	0
housing...	0	0
installp..	0	0
job.....	0	0
marital...	9	5
other.....	0	0
property..	0	0
purpose...	564	0
resident..	11	0
savings...	4	2
telephon..	19	0
good_bad..	0	0

\*\*\*\*\*

\*\*\*\*\* Running StatsModel - All Features with Validation \*\*\*\*\*

\*\*\*\*\*

Warning: Maximum number of iterations has been exceeded.

Current function value: 0.466458

Iterations: 35

Printing results summary...

### Logit Regression Results

```

=====
Dep. Variable:          good_bad    No. Observations:          700
Model:                  Logit       Df Residuals:              654
Method:                 MLE         Df Model:                45
Date:                   Thu, 11 Jun 2020    Pseudo R-squ.:          0.2533
Time:                   16:42:50           Log-Likelihood:         -326.52
converged:              False           LL-Null:                -437.29
Covariance Type:        nonrobust         LLR p-value:            3.670e-25
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	19.7128	1.95e+04	0.001	0.999	-3.82e+04	3.83e+04
age	0.0203	0.011	1.856	0.063	-0.001	0.042
amount	-9.472e-05	4.91e-05	-1.929	0.054	-0.000	1.5e-06
duration	-0.0246	0.011	-2.282	0.022	-0.046	-0.003
depends	-0.3219	0.305	-1.054	0.292	-0.921	0.277
foreign	1.4928	0.832	1.794	0.073	-0.138	3.124
telephon	0.2279	0.235	0.969	0.333	-0.233	0.689
checking1	-2.0323	0.276	-7.352	0.000	-2.574	-1.491
checking2	-1.4167	0.285	-4.979	0.000	-1.974	-0.859
checking3	-1.0766	0.425	-2.532	0.011	-1.910	-0.243
coapp1	-0.6176	0.462	-1.336	0.182	-1.524	0.289
coapp2	-1.3099	0.636	-2.060	0.039	-2.556	-0.064
employed1	0.2766	0.513	0.539	0.590	-0.728	1.282
employed2	-0.2890	0.357	-0.811	0.418	-0.988	0.410
employed3	0.1929	0.301	0.640	0.522	-0.398	0.784
employed4	0.4120	0.355	1.160	0.246	-0.284	1.108
existcr1	-15.3745	1.95e+04	-0.001	0.999	-3.83e+04	3.82e+04
existcr2	-15.7521	1.95e+04	-0.001	0.999	-3.83e+04	3.82e+04
existcr3	-15.9180	1.95e+04	-0.001	0.999	-3.83e+04	3.82e+04
history0	-1.8710	0.492	-3.803	0.000	-2.835	-0.907
history1	-1.6593	0.562	-2.951	0.003	-2.762	-0.557
history2	-1.0100	0.328	-3.081	0.002	-1.653	-0.367
history3	-0.8362	0.400	-2.088	0.037	-1.621	-0.051
housing1	-0.5449	0.568	-0.960	0.337	-1.658	0.568
housing2	-0.2042	0.544	-0.376	0.707	-1.270	0.861
installp1	0.8683	0.340	2.552	0.011	0.201	1.535
installp2	0.9388	0.276	3.396	0.001	0.397	1.481
installp3	0.5217	0.308	1.697	0.090	-0.081	1.124
job1	-0.7188	0.756	-0.951	0.342	-2.200	0.763
job2	-0.2276	0.410	-0.555	0.579	-1.031	0.576
job3	-0.3471	0.342	-1.016	0.310	-1.017	0.323
marital1	-0.5331	0.583	-0.914	0.361	-1.677	0.611
marital2	-0.5999	0.391	-1.535	0.125	-1.366	0.166
marital3	-0.0772	0.390	-0.198	0.843	-0.842	0.687
other1	-0.7518	0.293	-2.562	0.010	-1.327	-0.177
other2	-0.9835	0.444	-2.213	0.027	-1.855	-0.112
property1	0.3555	0.500	0.711	0.477	-0.625	1.336

property2	0.2710	0.489	0.554	0.579	-0.688	1.230
property3	0.3735	0.479	0.779	0.436	-0.566	1.313
resident1	0.4934	0.354	1.392	0.164	-0.201	1.188
resident2	-0.4395	0.266	-1.655	0.098	-0.960	0.081
resident3	-0.1930	0.318	-0.608	0.543	-0.815	0.429
savings1	-0.6255	0.296	-2.114	0.035	-1.206	-0.045
savings2	-0.4164	0.410	-1.017	0.309	-1.219	0.386
savings3	-0.4319	0.548	-0.789	0.430	-1.505	0.641
savings4	0.5030	0.750	0.671	0.502	-0.966	1.972

\*\*\*\*\* Training Model Metrics \*\*\*\*\*

Model Metrics

Observations.....	700
Accuracy.....	0.7814
Precision.....	0.8083
Sensitivity (Recall).....	0.8912
Specificity (Selectivity).. <td>0.5450</td>	0.5450
F1-Score.....	0.8478
MISC (Misclassification)...	21.9%
class 0.....	45.5%
class 1.....	10.9%

Confusion

Matrix	Class 0	Class 1
Class 0.....	121	101
Class 1.....	52	426

\*\*\*\*\* Validation Metrics \*\*\*\*\*

Model Metrics

Observations.....	300
Accuracy.....	0.7633
Precision.....	0.8186
Sensitivity (Recall).....	0.8739
Specificity (Selectivity).. <td>0.4487</td>	0.4487
F1-Score.....	0.8453
MISC (Misclassification)...	23.7%
class 0.....	55.1%
class 1.....	12.6%

Confusion

Matrix	Class 0	Class 1
Class 0.....	35	43
Class 1.....	28	194

\*\*\*\*\*

\*\*\*\*\*  
 \*\*\*\*\* Running StatsModel - Logistic with Stepwise \*\*\*\*\*  
 \*\*\*\*\*

```
/Users/edwardrainwater/opt/anaconda3/envs/Stat656MacOSX/lib/python3.7/site-  
packages/statsmodels/base/model.py:568: ConvergenceWarning: Maximum Likelihood  
optimization failed to converge. Check mle_retvals
```

```
"Check mle_retvals", ConvergenceWarning)
```

```
Add checking1 with p-value 1.54079e-15  
Add checking2 with p-value 1.00836e-14  
Add duration with p-value 1.04097e-09  
Add savings1 with p-value 0.000857473  
Add history1 with p-value 0.000921171  
Add employed4 with p-value 0.0028786  
Add savings2 with p-value 0.00643012  
Add marital3 with p-value 0.00758101  
Add history0 with p-value 0.00807971  
Add history2 with p-value 0.00770107  
Add property1 with p-value 0.0230649  
Add foreign with p-value 0.0379187  
Add resident2 with p-value 0.0280212  
Add housing2 with p-value 0.0380405  
Add employed2 with p-value 0.0557519  
Add other1 with p-value 0.0620334  
Add checking3 with p-value 0.0834477  
Add history3 with p-value 0.0981458  
Add existcr1 with p-value 0.0760241
```

```
Final selected attributes:
```

```
checking1  
checking2  
duration  
savings1  
history1  
employed4  
savings2  
marital3  
history0  
history2  
property1  
foreign  
resident2  
housing2  
employed2  
other1  
checking3  
history3  
existcr1
```

```
*****  
Optimization terminated successfully.
```

```
Current function value: 0.494159  
Iterations 7
```

```
*****  
***** Training Model *****
```

```
Target: good_bad  
Logit Regression Results
```

```
=====
```

```

Dep. Variable:          good_bad    No. Observations:          700
Model:                  Logit        Df Residuals:              680
Method:                  MLE         Df Model:                  19
Date:                   Thu, 11 Jun 2020    Pseudo R-squ.:            0.2090
Time:                   16:42:57          Log-Likelihood:           -345.91
converged:              True          LL-Null:                  -437.29
Covariance Type:        nonrobust        LLR p-value:              8.870e-29

```

	coef	std err	z	P> z	[0.025	0.975]
const	3.2929	0.426	7.737	0.000	2.459	4.127
checking1	-1.9458	0.259	-7.509	0.000	-2.454	-1.438
checking2	-1.1974	0.262	-4.562	0.000	-1.712	-0.683
duration	-0.0341	0.008	-4.200	0.000	-0.050	-0.018
savings1	-0.4201	0.241	-1.741	0.082	-0.893	0.053
history1	-1.8927	0.530	-3.569	0.000	-2.932	-0.853
employed4	0.2742	0.276	0.994	0.320	-0.266	0.815
savings2	-0.3846	0.357	-1.077	0.282	-1.084	0.315
marital3	0.2309	0.197	1.169	0.242	-0.156	0.618
history0	-1.7448	0.468	-3.730	0.000	-2.662	-0.828
history2	-0.9706	0.310	-3.129	0.002	-1.579	-0.363
property1	0.1400	0.220	0.637	0.524	-0.290	0.570
foreign	1.6184	0.809	2.001	0.045	0.033	3.203
resident2	-0.5057	0.210	-2.409	0.016	-0.917	-0.094
housing2	0.2714	0.211	1.288	0.198	-0.142	0.684
employed2	-0.5503	0.250	-2.200	0.028	-1.041	-0.060
other1	-0.5345	0.277	-1.932	0.053	-1.077	0.008
checking3	-0.8739	0.402	-2.175	0.030	-1.661	-0.087
history3	-0.6779	0.377	-1.797	0.072	-1.417	0.062
existcr1	0.4480	0.272	1.645	0.100	-0.086	0.982

```

*****
***** Training Model Metrics *****
*****

```

#### Model Metrics

```

Observations..... 700
Accuracy..... 0.7586
Precision..... 0.7932
Sensitivity (Recall)..... 0.8745
Specificity (Selectivity).. 0.5090
F1-Score..... 0.8318
MISC (Misclassification)... 24.1%
    class 0..... 49.1%
    class 1..... 12.6%

```

#### Confusion

```

    Matrix    Class 0    Class 1
Class 0.....    113     109
Class 1.....     60     418

```

```

*****

```

\*\*\*\*\* Validation Metrics \*\*\*\*\*

Model Metrics

Observations.....	300
Accuracy.....	0.7867
Precision.....	0.8238
Sensitivity (Recall).....	0.9054
Specificity (Selectivity).. <td>0.4487</td>	0.4487
F1-Score.....	0.8627
MISC (Misclassification)...	21.3%
class 0.....	55.1%
class 1.....	9.5%

Confusion

Matrix	Class 0	Class 1
Class 0.....	35	43
Class 1.....	21	201

\*\*\*\*\*Regularization Logistic Regression\*\*\*\*\*

Logistic Regression Model using C= 0.0001

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	99	99
Mean Absolute Error....	0.4161	0.3915
Avg Squared Error.....	0.2064	0.1835
Accuracy.....	0.6771	0.7500
Precision.....	0.6886	0.7492
Recall (Sensitivity).....	0.9623	0.9955
F1-score.....	0.8028	0.8549
Total Misclassifications...	226	75
MISC (Misclassification)...	32.3%	25.0%
class 0.....	93.7%	94.9%
class 1.....	3.8%	0.5%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	14	208
Class 1.....	18	460

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	4	74
Class 1.....	1	221

Logistic Regression Model using C= 0.01



Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	50	50
Mean Absolute Error....	0.3977	0.3765
Avg Squared Error.....	0.1926	0.1753
Accuracy.....	0.7114	0.7700
Precision.....	0.7184	0.7762
Recall (Sensitivity).....	0.9498	0.9685
F1-score.....	0.8180	0.8617
Total Misclassifications...	202	69
MISC (Misclassification)...	28.9%	23.0%
class 0.....	80.2%	79.5%
class 1.....	5.0%	3.2%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	44	178
Class 1.....	24	454

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	16	62
Class 1.....	7	215

Logistic Regression Model using C= 0.1

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	126	126
Mean Absolute Error....	0.3548	0.3396
Avg Squared Error.....	0.1676	0.1587
Accuracy.....	0.7557	0.7900
Precision.....	0.7707	0.8118
Recall (Sensitivity).....	0.9142	0.9324
F1-score.....	0.8364	0.8679
Total Misclassifications...	171	63
MISC (Misclassification)...	24.4%	21.0%
class 0.....	58.6%	61.5%
class 1.....	8.6%	6.8%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	92	130
Class 1.....	41	437

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	30	48
Class 1.....	15	207

Logistic Regression Model using C= 1.0

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	134	134
Mean Absolute Error....	0.3282	0.3197
Avg Squared Error.....	0.1614	0.1588
Accuracy.....	0.7571	0.7433
Precision.....	0.7906	0.8139
Recall (Sensitivity).....	0.8766	0.8468
F1-score.....	0.8313	0.8300
Total Misclassifications...	170	77
MISC (Misclassification)...	24.3%	25.7%
class 0.....	50.0%	55.1%
class 1.....	12.3%	15.3%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	111	111
Class 1.....	59	419

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	35	43
Class 1.....	34	188

Logistic Regression Model using C= 5.0

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	848	848
Mean Absolute Error....	0.3118	0.3135
Avg Squared Error.....	0.1552	0.1649
Accuracy.....	0.7757	0.7533
Precision.....	0.8104	0.8162
Recall (Sensitivity).....	0.8766	0.8604
F1-score.....	0.8422	0.8377
Total Misclassifications...	157	74
MISC (Misclassification)...	22.4%	24.7%

class 0.....	44.1%	55.1%
class 1.....	12.3%	14.0%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	124	98
Class 1.....	59	419

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	35	43
Class 1.....	31	191

Logistic Regression Model using C= 10.0

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	154	154
Mean Absolute Error....	0.3232	0.3161
Avg Squared Error.....	0.1613	0.1600
Accuracy.....	0.7571	0.7433
Precision.....	0.7928	0.8194
Recall (Sensitivity).....	0.8724	0.8378
F1-score.....	0.8307	0.8285
Total Misclassifications...	170	77
MISC (Misclassification)...	24.3%	25.7%
class 0.....	49.1%	52.6%
class 1.....	12.8%	16.2%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	113	109
Class 1.....	61	417

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	37	41
Class 1.....	36	186

Logistic Regression Model using C= 50.0

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	691	691

Mean Absolute Error....	0.3135	0.3150
Avg Squared Error.....	0.1558	0.1649
Accuracy.....	0.7757	0.7600
Precision.....	0.8046	0.8178
Recall (Sensitivity).....	0.8870	0.8694
F1-score.....	0.8438	0.8428
Total Misclassifications...	157	72
MISC (Misclassification)...	22.4%	24.0%
class 0.....	46.4%	55.1%
class 1.....	11.3%	13.1%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	119	103
Class 1.....	54	424

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	35	43
Class 1.....	29	193

Logistic Regression Model using C= inf

Model Metrics.....	Training	Validation
Observations.....	700	300
Coefficients.....	46	46
DF Error.....	654	254
Iterations.....	484	484
Mean Absolute Error....	0.3108	0.3122
Avg Squared Error.....	0.1558	0.1650
Accuracy.....	0.7757	0.7667
Precision.....	0.8081	0.8248
Recall (Sensitivity).....	0.8808	0.8694
F1-score.....	0.8428	0.8465
Total Misclassifications...	157	70
MISC (Misclassification)...	22.4%	23.3%
class 0.....	45.0%	52.6%
class 1.....	11.9%	13.1%

Training	Class	Class
Confusion Matrix	0	1
Class 0.....	122	100
Class 1.....	57	421

Validation	Class	Class
Confusion Matrix	0	1
Class 0.....	37	41
Class 1.....	29	193

```

** Cross-Validation for Regularization Logistic Regression **
0.000100..    0.8171    0.0144
0.010000..    0.8285    0.0177
0.100000..    0.8324    0.0341
1.000000..    0.8260    0.0274
5.000000..    0.8335    0.0288
10.000000..   0.8266    0.0278
50.000000..   0.8275    0.0271
inf.....    0.8315    0.0321

```

### 3 Appendix – Python Code Listing

```

"""
Created 09 JUN 2020

@author: el-rainwater, Rainwater Center for Neolithic Computing
"""

import pandas as pd
import numpy as np
from AdvancedAnalytics.ReplaceImputeEncode import ReplaceImputeEncode, DT
from AdvancedAnalytics.Regression import logreg, stepwise
import statsmodels.api as sm
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression

filepath = '/Users/edwardrainwater/OneDrive - Texas A&M University/' \
'Summer-2020/STAT 656 Applied Analytics/hw-03/'
file = 'credithistory.xlsx'

def run_all_features(encoded_df,target):
    print("\n" + ("*" * 78))
    print("*" * 13 + "    Running StatsModel - All Features with Validation    " + "*" * 13)
    print(("*" * 78))
    y = encoded_df[target].astype(int)
    X = encoded_df.drop(target, axis=1)
    Xt, Xv, yt, yv = train_test_split(X, y, train_size=0.7, random_state=12345)
    Xtc = sm.add_constant(Xt)
    Xvc = sm.add_constant(Xv)
    model = sm.Logit(yt, Xtc)
    results = model.fit()
    print('Printing results summary...')
    print(results.summary())

    print("***** Training Model Metrics *****")
    mat = results.pred_table(threshold=0.5)
    logreg.display_confusion(mat)

    print("\n***** Validation Metrics *****")
    predv = results.predict(Xvc)
    sv = np.where(predv<0.5, 0, 1)
    logreg.display_confusion(pd.crosstab(sv, yv))
    print("*****\n")

```

```

return()

def run_stepwise(df_encoded,target):
    print("\n" + ("*" * 78))
    print("*" * 15 + "    Running StatsModel - Logistic with Stepwise    " + "*" * 15)
    print(("*" * 78))

    # Set up stepwise feature selection
    df_encoded[target] = df_encoded[target].astype(int) # Do this to make target int
    y = df_encoded[target]
    sw = stepwise(df_encoded, target, reg='logistic', verbose=True)

    selected = sw.fit_transform()

    print('\nFinal selected attributes:')
    print(*selected,sep='\n')
    print(("*" * 78))

    # Split the model 70/30 for training/validation

    X_train, X_validate, y_train, y_validate = \
        train_test_split(df_encoded[selected], y, train_size=0.7,
                        random_state = 12345)

    Xc_train    = sm.add_constant(X_train)
    Xc_validate = sm.add_constant(X_validate)

    model      = sm.Logit(y_train, Xc_train)
    results    = model.fit()
    print("\n" + ("*" * 78))
    print("*" * 31 + " Training Model " + "*" * 31)
    print("                                Target: " + target)
    print(results.summary())

    print("\n" + ("*" * 78))
    print("*" * 27 + " Training Model Metrics " + "*" * 27)
    print(("*" * 80))

    mat = results.pred_table(threshold=0.5)
    logreg.display_confusion(mat)

    print("\n" + ("*" * 78))
    print("*" * 29 + " Validation Metrics " + "*" * 29)
    print(("*" * 78))

    predv = results.predict(Xc_validate)
    sv     = np.where(predv<0.5, 0, 1)
    logreg.display_confusion(pd.crosstab(sv, y_validate))
    print(("*" * 78))

    return()

def run_reglr_logistic_regression(encoded_df, target):

```

```

print("*****Regularization Logistic Regression*****")
y = encoded_df[target].astype(int)
X = encoded_df.drop(target, axis=1)
X_train, X_validate, y_train, y_validate = \
train_test_split(X, y, train_size=0.7, random_state = 12345)

C_list = [1e-4, 1e-2, 1e-1, 1.0, 5.0, 10.0, 50.0, np.inf]
for c in C_list:
    lr = LogisticRegression(C=c, tol=1e-4, solver='lbfgs', max_iter=5000)
    lr = lr.fit(X_train, y_train)
    print("\nLogistic Regression Model using C=", c)
    logreg.display_split_metrics(lr, X_train, y_train, X_validate,
                                y_validate, target_names=['Bad', 'Good'])

print("\n** Cross-Validation for Regularization Logistic Regression **")
for c in C_list:
    lr = LogisticRegression(C=c, tol=1e-4, solver='lbfgs', max_iter=5000)
    lrc = cross_val_score(lr, X, y, cv=10, scoring='f1', n_jobs=3)
    mean = lrc.mean()
    std = lrc.std()
    print("{:.<10.6f}{:}>10.4f}{:}>10.4f}".format(c, mean, std))
return()

```

```

def main(): #####

```

```

df = pd.read_excel(filepath + file)
print(df.head(), df.shape)
print(df.dtypes)

attribute_map = {
    'age':[DT.Interval, (19,120)],
    'amount':[DT.interval, (0,20000)],
    'checking':[DT.nominal, (list(range(1,5)))],
    'coapp':[DT.nominal, (1,2,3)],
    'depends':[DT.Binary, (1,2)],
    'duration':[DT.interval, (1,72)],
    'employed':[DT.Nominal, (list(range(1,6)))],
    'existcr':[DT.Nominal, (1,2,3,4)],
    'foreign':[DT.Binary, (1,2)],
    'history':[DT.nominal, (0,1,2,3,4)],
    'housing':[DT.Nominal, (1,2,3)],
    'installp':[DT.nominal, (1,2,3,4)],
    'job':[DT.nominal, (1,2,3,4)],
    'marital':[DT.Nominal, (1,2,3,4)],
    'other':[DT.Nominal, (1,2,3)],
    'property':[DT.nominal, (1,2,3,4)],
    'purpose':[DT.Ignore, ('0', '1', '2', '3', '4', '5', \
                           '6', '8', '9', 'X') ],
    'resident':[DT.Nominal, (1,2,3,4)],
    'savings':[DT.nominal,(1,2,3,4,5)],
    'telephon':[DT.Binary, (1,2)],
    'good_bad':[DT.Binary , ('bad', 'good') ]
}

```

```
}
```

```
target = 'good_bad'
```

```
# One-hot encode and impute missing values
```

```
rie = ReplaceImputeEncode(data_map=attribute_map, nominal_encoding='one-hot',  
                           binary_encoding='one-hot', no_impute=[target],  
                           interval_scale=None, drop=True,  
                           display=True)
```

```
df_encoded = rie.fit_transform(df).dropna() #drop rows with missing values
```

```
run_all_features(df_encoded, target)
```

```
run_stepwise(df_encoded, target)
```

```
run_reglr_logistic_regression(df_encoded, target)
```

```
if __name__ == '__main__':  
    main()
```